

D206 Performance Assessment By Krista Moik

```
In [1]: #import all libraries
import pandas as pd
import math
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats
import seaborn as sns
```

```
In [2]: #import CSV
df = pd.read_csv('C:/Users/Kmoik WGU/Desktop/medical_raw_data.csv')
```

```
In [3]: #view the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 53 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        10000 non-null   int64  
 1   CaseOrder         10000 non-null   int64  
 2   Customer_id       10000 non-null   object  
 3   Interaction       10000 non-null   object  
 4   UID               10000 non-null   object  
 5   City              10000 non-null   object  
 6   State             10000 non-null   object  
 7   County            10000 non-null   object  
 8   Zip               10000 non-null   int64  
 9   Lat               10000 non-null   float64 
 10  Lng               10000 non-null   float64 
 11  Population        10000 non-null   int64  
 12  Area              10000 non-null   object  
 13  Timezone          10000 non-null   object  
 14  Job               10000 non-null   object  
 15  Children          7412 non-null   float64 
 16  Age               7586 non-null   float64 
 17  Education         10000 non-null   object  
 18  Employment        10000 non-null   object  
 19  Income             7536 non-null   float64 
 20  Marital            10000 non-null   object  
 21  Gender             10000 non-null   object  
 22  ReAdmis            10000 non-null   object  
 23  VitD_levels        10000 non-null   float64 
 24  Doc_visits         10000 non-null   int64  
 25  Full_meals_eaten   10000 non-null   int64  
 26  VitD_supp          10000 non-null   int64  
 27  Soft_drink          7533 non-null   object  
 28  Initial_admin       10000 non-null   object  
 29  HighBlood           10000 non-null   object  
 30  Stroke              10000 non-null   object  
 31  Complication_risk   10000 non-null   object  
 32  Overweight          9018 non-null   float64 
 33  Arthritis            10000 non-null   object  
 34  Diabetes             10000 non-null   object  
 35  Hyperlipidemia       10000 non-null   object  
 36  BackPain             10000 non-null   object  
 37  Anxiety              9016 non-null   float64 
 38  Allergic_rhinitis    10000 non-null   object  
 39  Reflux_esophagitis    10000 non-null   object  
 40  Asthma               10000 non-null   object  
 41  Services              10000 non-null   object  
 42  Initial_days          8944 non-null   float64 
 43  TotalCharge           10000 non-null   float64 
 44  Additional_charges     10000 non-null   float64 
 45  Item1                10000 non-null   int64  
 46  Item2                10000 non-null   int64  
 47  Item3                10000 non-null   int64  
 48  Item4                10000 non-null   int64  
 49  Item5                10000 non-null   int64  
 50  Item6                10000 non-null   int64  
 51  Item7                10000 non-null   int64  
 52  Item8                10000 non-null   int64  
dtypes: float64(11), int64(15), object(27)
memory usage: 4.0+ MB
```

This view of the data shows us we have 53 columns and 10,000 rows, which is correct. This view also shows that we do have NULL values that we will need to look at.

```
In [4]: #Viewing the number of duplicates in the data  
print(df.duplicated().value_counts())
```

```
False    10000  
Name: count, dtype: int64
```

The output of the sum of duplicate shows we do not have any duplicate values.

```
In [5]: # We know from our df.info() view there are NULL values. We will search for the sum of  
df.isnull().sum()
```

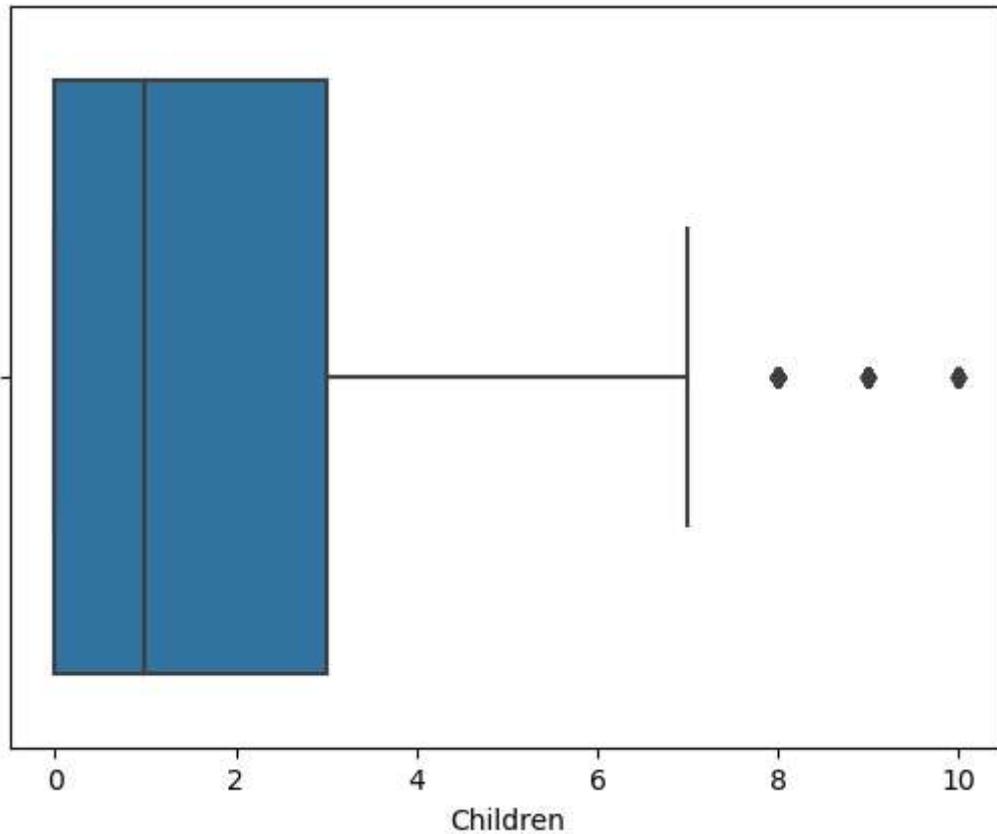
```
Out[5]: Unnamed: 0      0
CaseOrder          0
Customer_id        0
Interaction        0
UID                0
City               0
State              0
County             0
Zip                0
Lat                0
Lng                0
Population         0
Area               0
Timezone           0
Job                0
Children           2588
Age                2414
Education          0
Employment         0
Income              2464
Marital             0
Gender              0
ReAdmis             0
VitD_levels         0
Doc_visits          0
Full_meals_eaten    0
VitD_supp            0
Soft_drink           2467
Initial_admin        0
HighBlood            0
Stroke              0
Complication_risk    0
Overweight           982
Arthritis            0
Diabetes             0
Hyperlipidemia       0
BackPain             0
Anxiety              984
Allergic_rhinitis    0
Reflux_esophagitis   0
Asthma               0
Services              0
Initial_days          1056
TotalCharge           0
Additional_charges    0
Item1                0
Item2                0
Item3                0
Item4                0
Item5                0
Item6                0
Item7                0
Item8                0
dtype: int64
```

The output shows that there are 2588 NULL values for children, 2414 NULL values for Age, 2464 NULL values for Income, 2467 NULL values for soft_drink, 982 NULL values for overweight, 984 NULL values for anxiety, and 1056 NULL values for Initial_days.

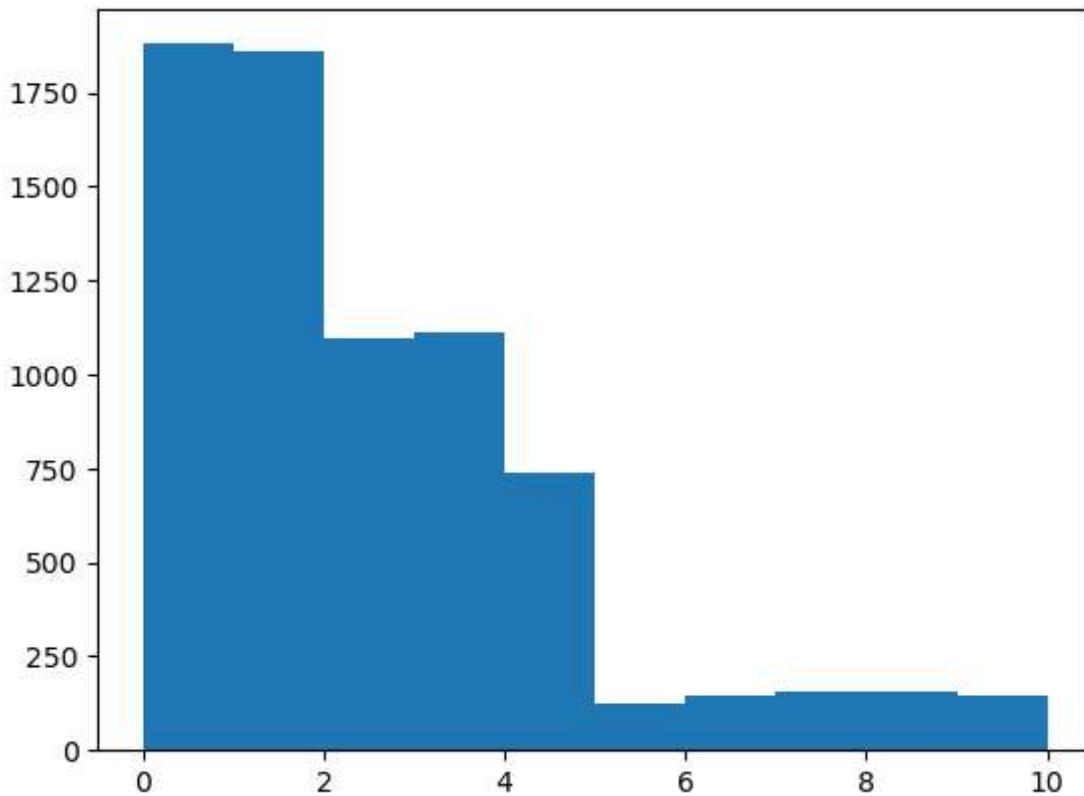
We will now examine the NULL values of the Quantitative columns - Children, Age, Income, and Initial_days using Boxplots and Histograms

```
In [6]: #Boxplot for Children  
sns.boxplot(df, x='Children')
```

```
Out[6]: <Axes: xlabel='Children'>
```



```
In [7]: #Histogram for Children  
plt.hist(df['Children'])  
plt.show()
```



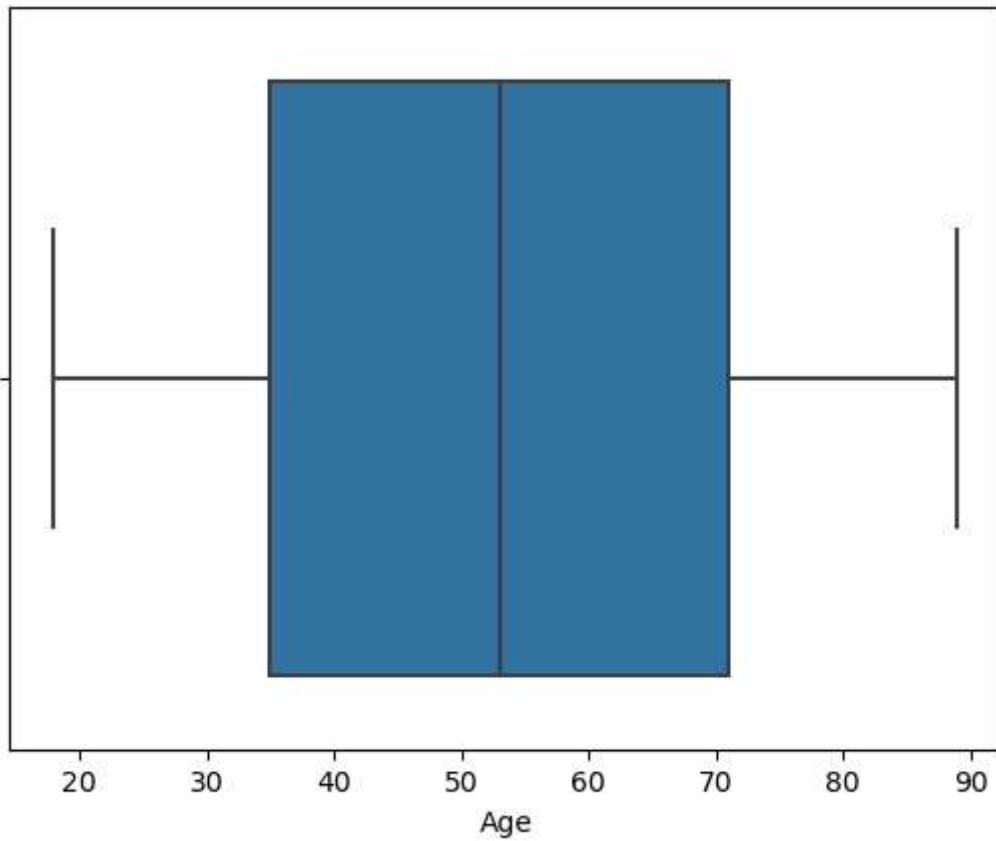
```
In [8]: #obtain statistical information regarding the column  
df.Children.describe()
```

```
Out[8]: count    7412.000000  
mean      2.098219  
std       2.155427  
min       0.000000  
25%      0.000000  
50%      1.000000  
75%      3.000000  
max      10.000000  
Name: Children, dtype: float64
```

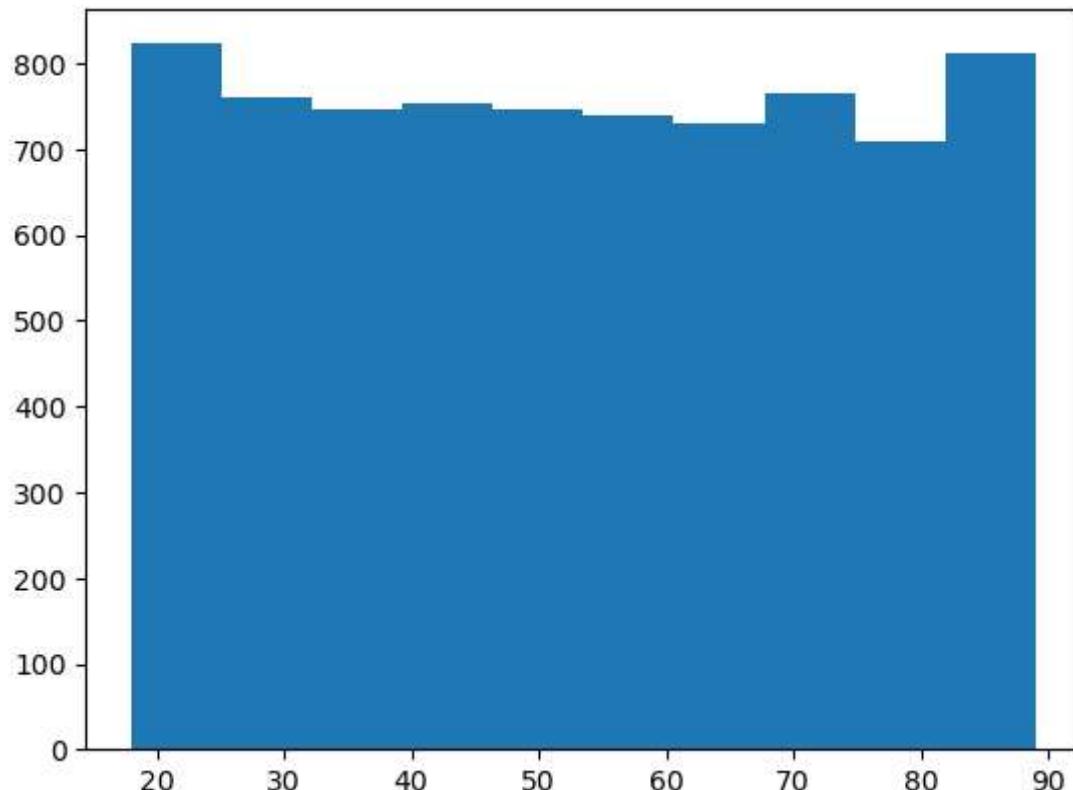
There appear to be 3 outliers in this column as well as NULL values. This histogram shows a right skew.

```
In [9]: #Boxplot for Age  
sns.boxplot(df, x='Age')
```

```
Out[9]: <Axes: xlabel='Age'>
```



```
In [10]: #Histogram for Age  
plt.hist(df['Age'])  
plt.show()
```



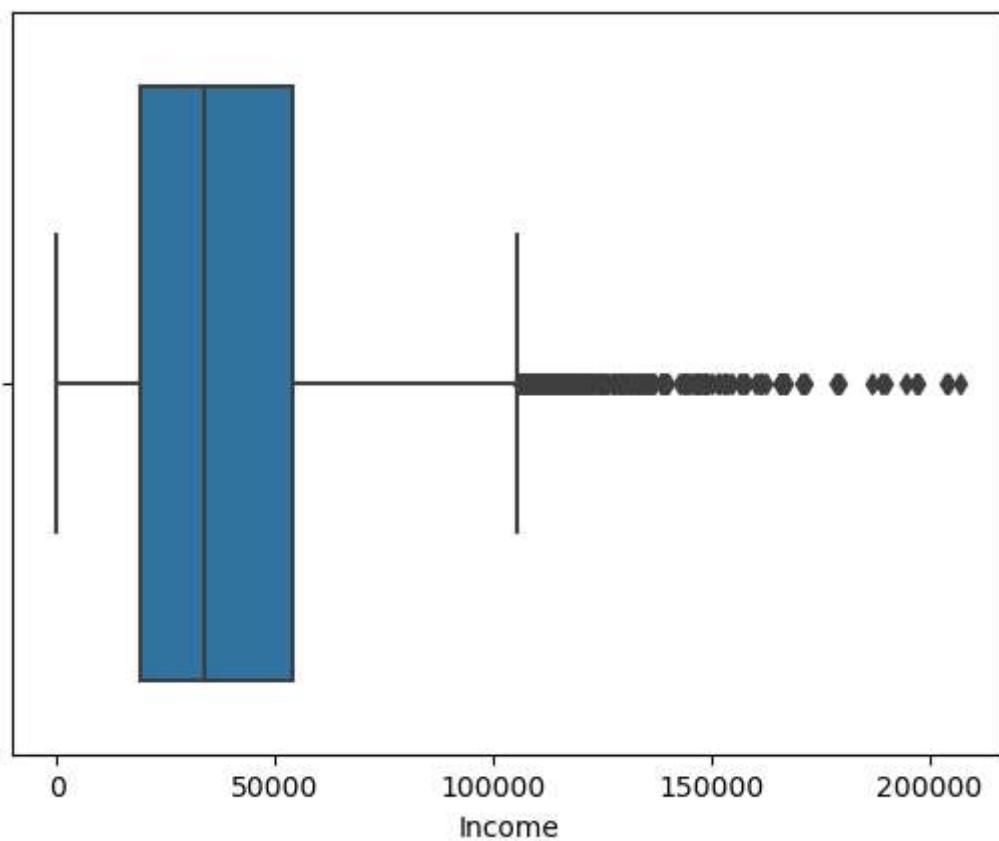
```
In [11]: #obtain statistical information regarding the column  
df.Age.describe()
```

```
Out[11]: count    7586.000000
          mean     53.295676
          std      20.659182
          min      18.000000
          25%     35.000000
          50%     53.000000
          75%     71.000000
          max     89.000000
Name: Age, dtype: float64
```

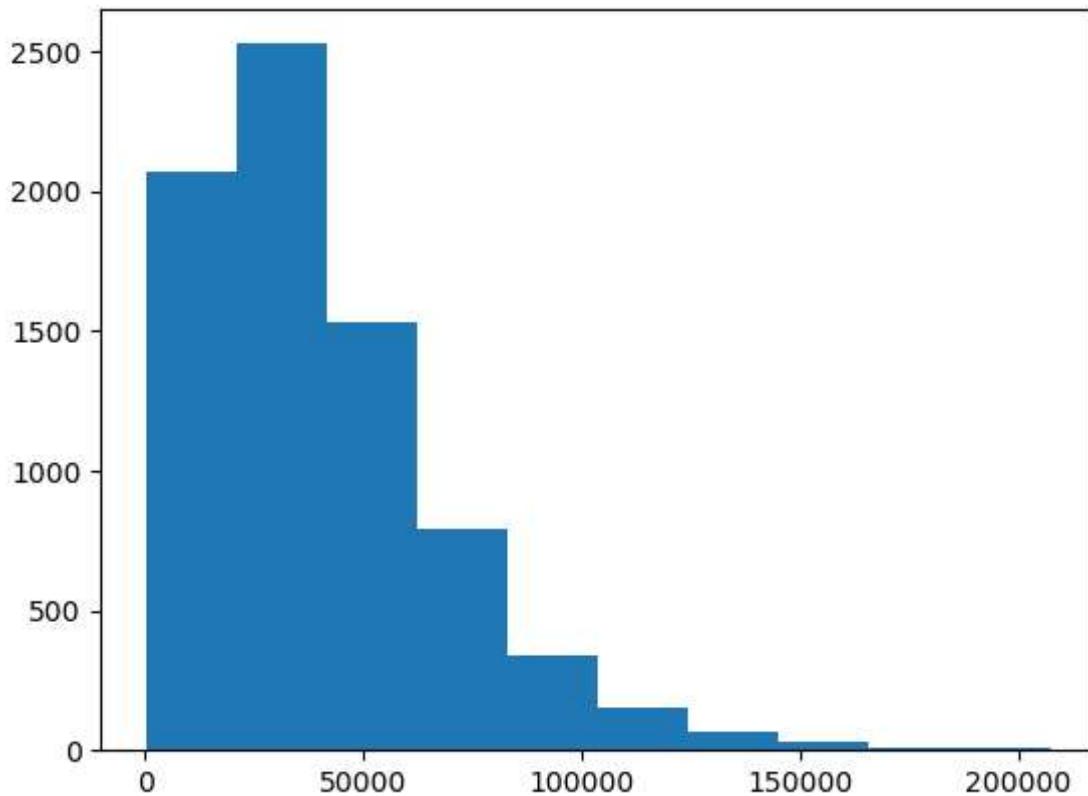
There do not appear to be any outliers in this column. This histogram shows a mostly uniform distribution.

```
In [12]: #Boxplot for Income
sns.boxplot(df, x='Income')
```

```
Out[12]: <Axes: xlabel='Income'>
```



```
In [13]: #Histogram for Income
plt.hist(df['Income'])
plt.show()
```



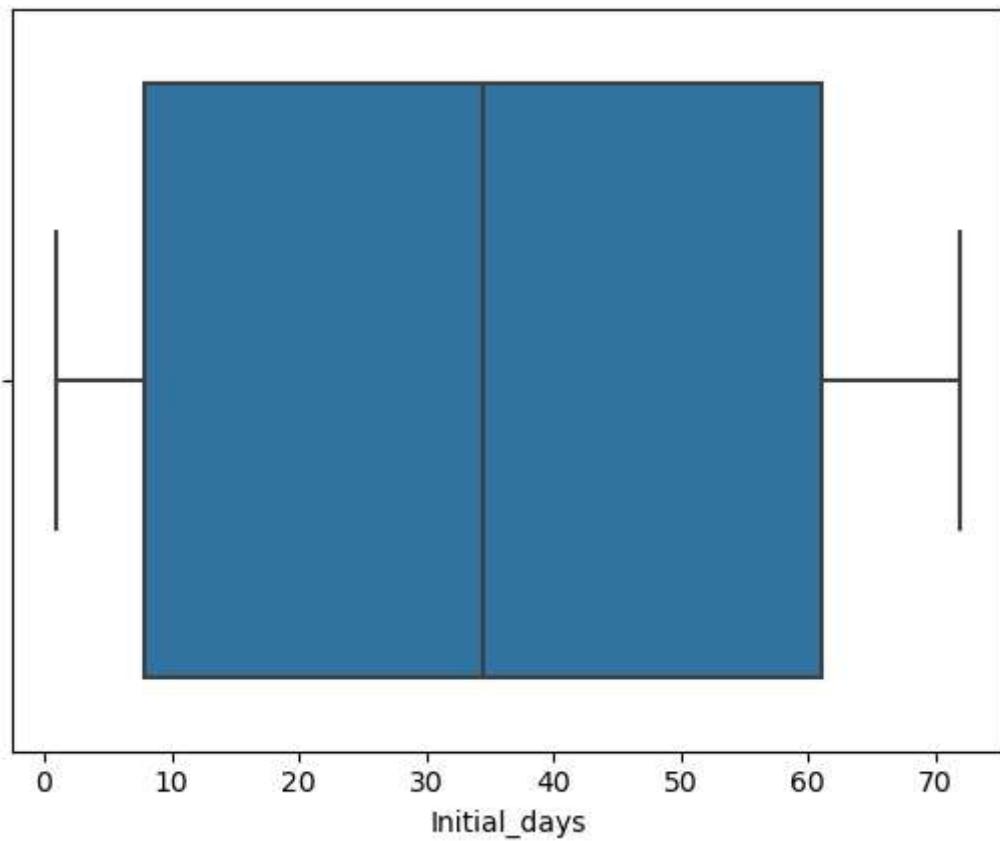
```
In [14]: #Obtain statistical information for the column  
df.Income.describe()
```

```
Out[14]: count    7536.000000  
mean     40484.438268  
std      28664.861050  
min      154.080000  
25%     19450.792500  
50%     33942.280000  
75%     54075.235000  
max     207249.130000  
Name: Income, dtype: float64
```

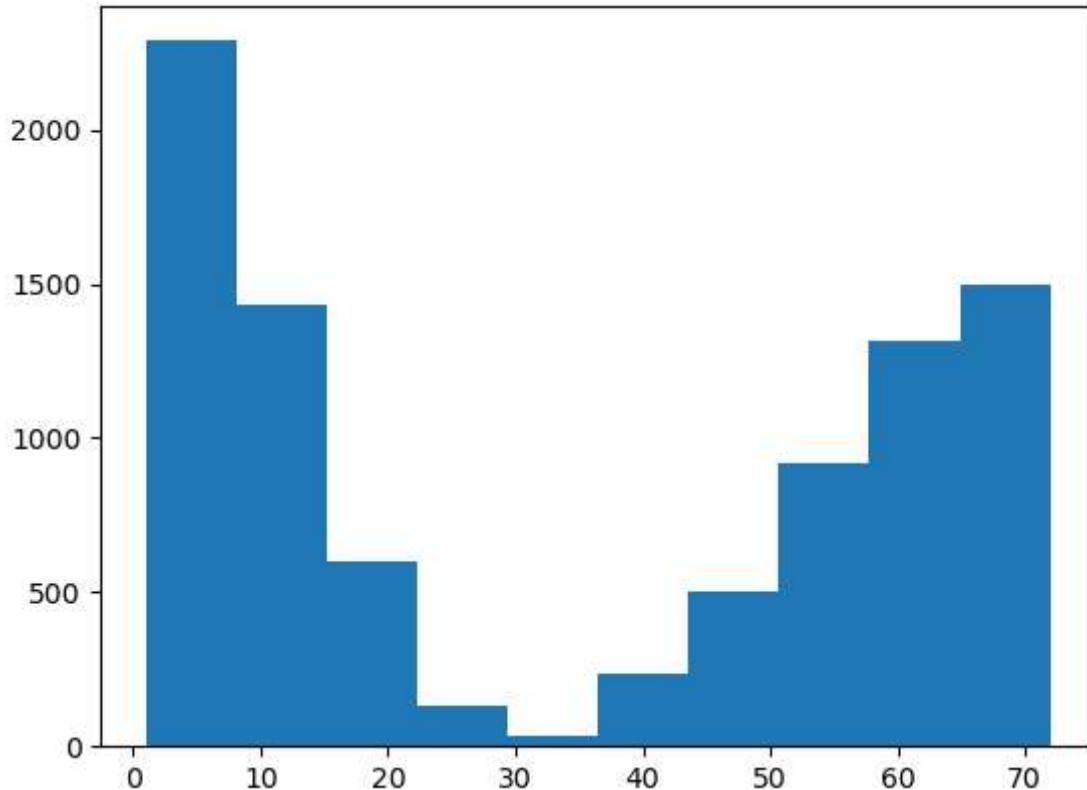
There appear to be many outliers in addition to NULL values in this column. This Histogram shows a right skew.

```
In [15]: #Boxplot for Initial Days  
sns.boxplot(df, x='Initial_days')
```

```
Out[15]: <Axes: xlabel='Initial_days'>
```



```
In [16]: #Histogram of Initial days  
plt.hist(df['Initial_days'])  
plt.show()
```



```
In [17]: #Obtain statistical information about this column  
df.Initial_days.describe()
```

```
Out[17]: count    8944.000000
          mean     34.432082
          std      26.287050
          min      1.001981
          25%      7.911709
          50%      34.446941
          75%      61.124654
          max      71.981486
Name: Initial_days, dtype: float64
```

There do not appear to be any outliers in this column. This histogram shows a bimodal skew.

We will now impute the NULL values before handling the outliers. We will use the mean for the normal distribution of Age, and the median for the right skew and bimodel columns.

```
In [18]: #imputing null values median and mean with code found in WGU Course Materials
df['Children'].fillna(df['Children'].median(), inplace=True)
df['Income'].fillna(df['Income'].median(), inplace=True)
df['Initial_days'].fillna(df['Initial_days'].median(), inplace=True)
df['Age'].fillna(df['Age'].mean(), inplace=True)
```

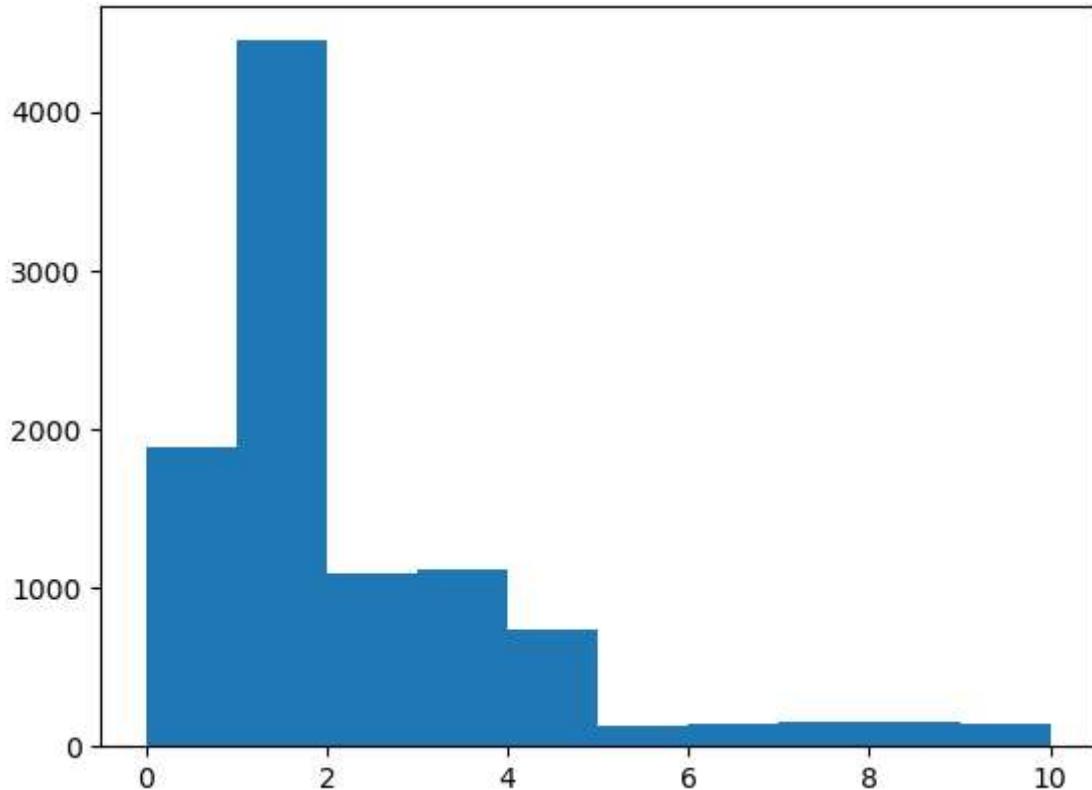
```
In [19]: #verify our imputation of NULL values in the quantitative columns
df.isnull().sum()
```

```
Out[19]: Unnamed: 0      0
CaseOrder          0
Customer_id        0
Interaction        0
UID                0
City               0
State              0
County             0
Zip                0
Lat                0
Lng                0
Population         0
Area               0
Timezone           0
Job                0
Children           0
Age                0
Education          0
Employment         0
Income              0
Marital             0
Gender              0
ReAdmis            0
VitD_levels        0
Doc_visits         0
Full_meals_eaten   0
VitD_supp          0
Soft_drink          2467
Initial_admin       0
HighBlood           0
Stroke              0
Complication_risk  0
Overweight          982
Arthritis           0
Diabetes            0
Hyperlipidemia     0
BackPain            0
Anxiety             984
Allergic_rhinitis  0
Reflux_esophagitis 0
Asthma              0
Services            0
Initial_days        0
TotalCharge         0
Additional_charges 0
Item1               0
Item2               0
Item3               0
Item4               0
Item5               0
Item6               0
Item7               0
Item8               0
dtype: int64
```

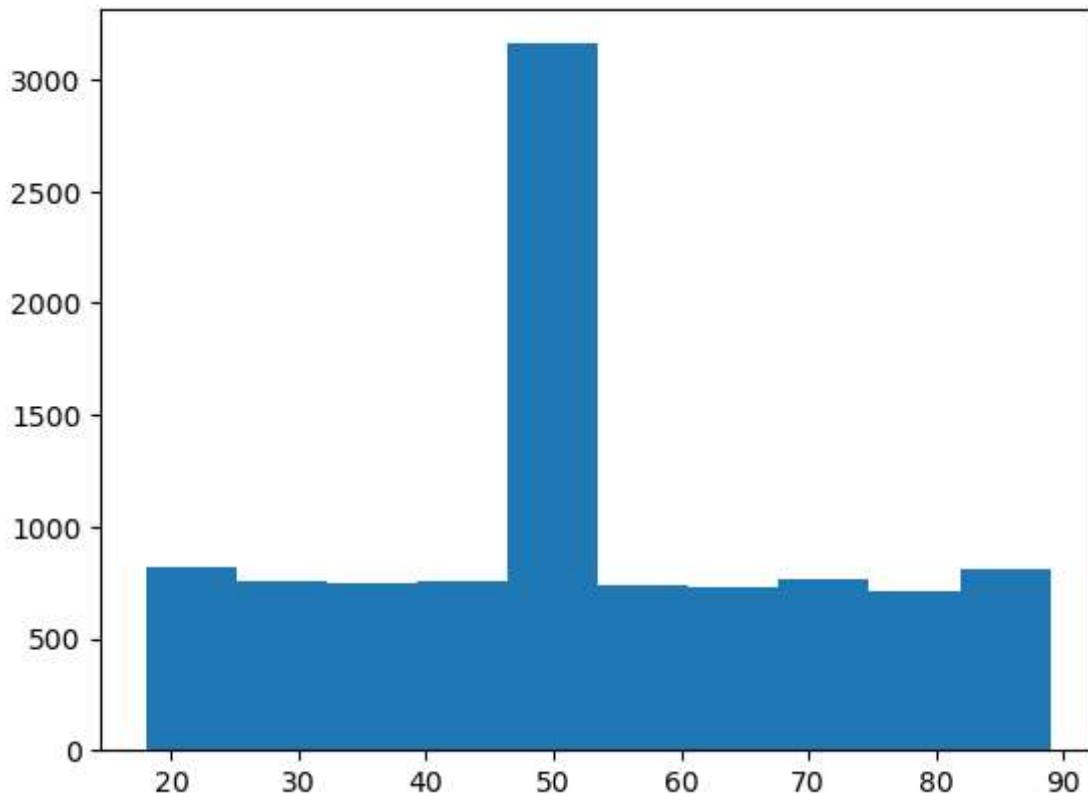
This shows our imputation of the quantitative columns was successful. The remaining NULLS are in qualitative columns.

Visual Confirmation Imputation of the NULLS was successful

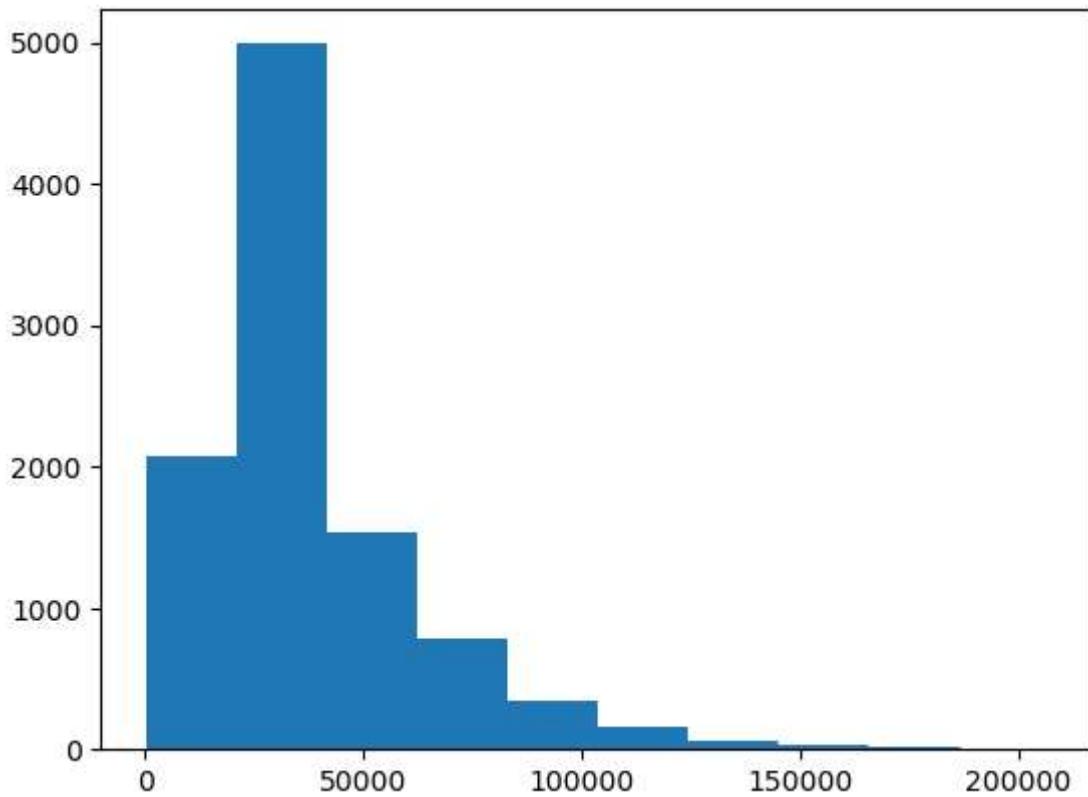
```
In [20]: #Children column  
plt.hist(df['Children'])  
plt.show()
```



```
In [21]: #Age column  
plt.hist(df['Age'])  
plt.show()
```

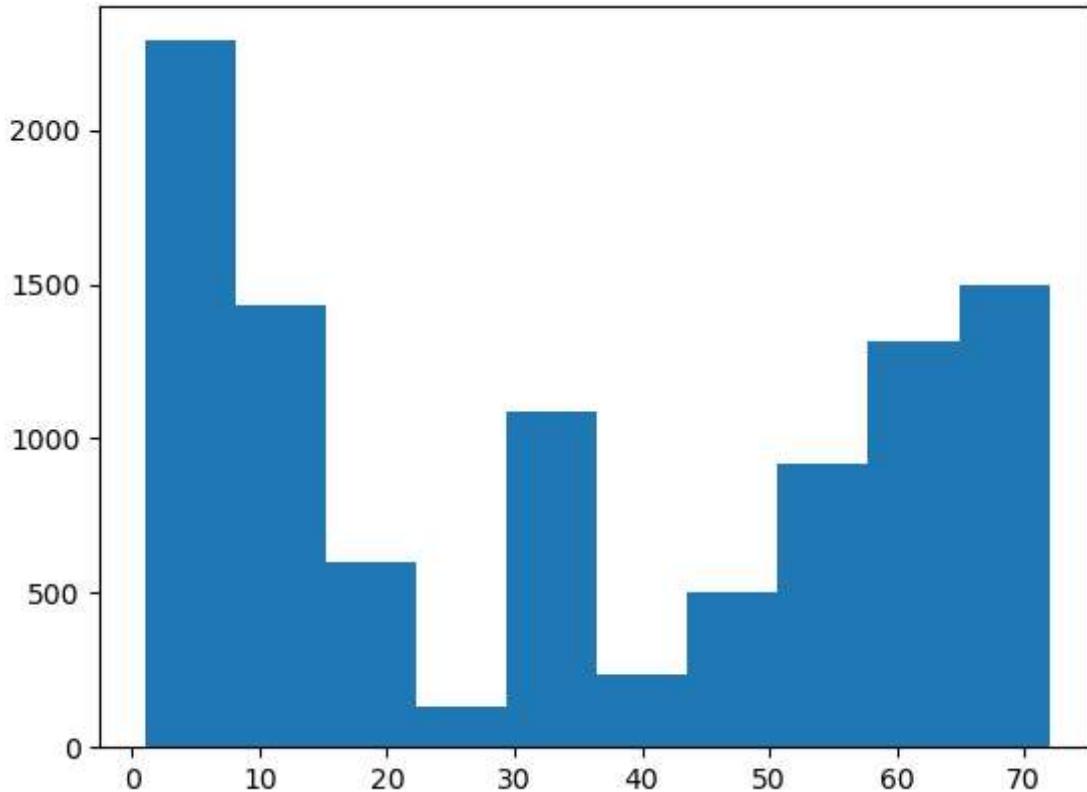


```
In [22]: #Income Column  
plt.hist(df['Income'])  
plt.show()
```



```
In [23]: #Initial_days column  
plt.hist(df['Initial_days'])
```

```
plt.show()
```



Imputing NULLS in qualitative columns

In [24]: # Find value counts excluding NULL values in Qualitative column Soft Drinks using code
df['Soft_drink'].value_counts(normalize=True, sort=True, dropna=True)

Out[24]:

Soft_drink	proportion
No	0.741935
Yes	0.258065

Name: proportion, dtype: float64

In [25]: #Find value counts excluding NULL values in Qualitative column Overweight using code f
df['Overweight'].value_counts(normalize=True, sort=True, dropna=True)

Out[25]:

Overweight	proportion
1.0	0.709137
0.0	0.290863

Name: proportion, dtype: float64

In [26]: #Find value counts excluded NULL values in Qualitative column Anxiety using code from
df['Anxiety'].value_counts(normalize=True, sort=True, dropna=True)

Out[26]:

Anxiety	proportion
0.0	0.677684
1.0	0.322316

Name: proportion, dtype: float64

Now that we have the count of the yes/no 0/1 responses and because these values are categorical, we will replace the NULL values with the most common response in each column.

In [27]: # 0 is the most common response for Soft drinks so we will impute the NULLS to No as t
df['Soft_drink'].fillna('No', inplace=True)

```
In [28]: # Using Ordinal encoding - Replicate the variable in preparation for replacing its cat
df['Soft_drink_numeric'] = df['Soft_drink']
```

```
In [29]: # Set up a dictionary specifically for converting the categorical values to numeric va
dict_soft_drink = {'Soft_drink_numeric': {'No': 0, 'Yes': 1}}
```

```
In [30]: # Use the dictionary to replace the variable's values. The replace() function will rep
df.replace(dict_soft_drink, inplace=True)
```

```
In [31]: #check to see if our dictionary worked for soft_drink_numeric using code from WGU cour
df['Soft_drink_numeric'].value_counts(normalize=True, sort=True, dropna=True)
```

```
Out[31]: Soft_drink_numeric
0      0.8056
1      0.1944
Name: proportion, dtype: float64
```

We can see our Ordinal encoding was successful. We will now continue replacing the remaining NULLS

```
In [32]: # 1 is the most common response for Overweight so we will impute the NULLS to 1 using
df['Overweight'].fillna(1, inplace=True)
```

```
In [33]: # 0 is the most common response for Anxiety so we will impute the NULLS to 0 using coa
df['Anxiety'].fillna(0, inplace=True)
```

```
In [34]: # View dataset to verify all NULLS have been addressed
df.info()
```

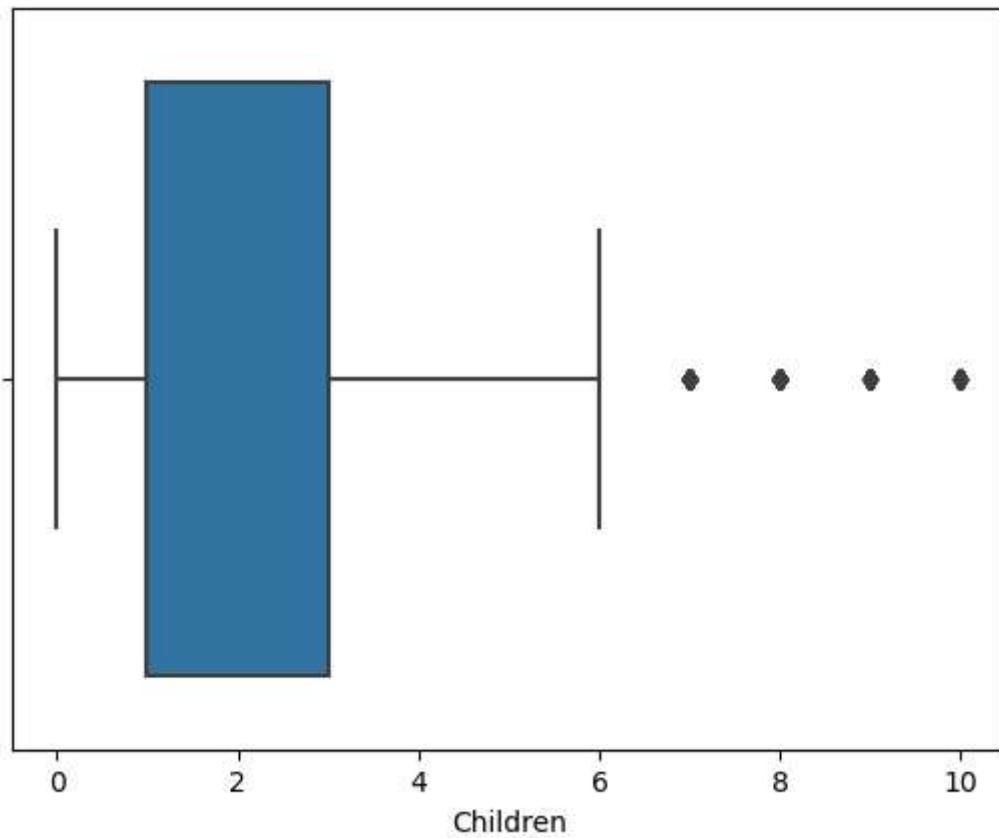
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 54 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        10000 non-null   int64  
 1   CaseOrder         10000 non-null   int64  
 2   Customer_id       10000 non-null   object  
 3   Interaction       10000 non-null   object  
 4   UID               10000 non-null   object  
 5   City              10000 non-null   object  
 6   State             10000 non-null   object  
 7   County            10000 non-null   object  
 8   Zip               10000 non-null   int64  
 9   Lat               10000 non-null   float64 
 10  Lng               10000 non-null   float64 
 11  Population        10000 non-null   int64  
 12  Area              10000 non-null   object  
 13  Timezone          10000 non-null   object  
 14  Job               10000 non-null   object  
 15  Children          10000 non-null   float64 
 16  Age               10000 non-null   float64 
 17  Education         10000 non-null   object  
 18  Employment        10000 non-null   object  
 19  Income             10000 non-null   float64 
 20  Marital            10000 non-null   object  
 21  Gender             10000 non-null   object  
 22  ReAdmis            10000 non-null   object  
 23  VitD_levels        10000 non-null   float64 
 24  Doc_visits         10000 non-null   int64  
 25  Full_meals_eaten   10000 non-null   int64  
 26  VitD_supp          10000 non-null   int64  
 27  Soft_drink          10000 non-null   object  
 28  Initial_admin       10000 non-null   object  
 29  HighBlood           10000 non-null   object  
 30  Stroke              10000 non-null   object  
 31  Complication_risk   10000 non-null   object  
 32  Overweight          10000 non-null   float64 
 33  Arthritis            10000 non-null   object  
 34  Diabetes             10000 non-null   object  
 35  Hyperlipidemia       10000 non-null   object  
 36  BackPain             10000 non-null   object  
 37  Anxiety              10000 non-null   float64 
 38  Allergic_rhinitis    10000 non-null   object  
 39  Reflux_esophagitis   10000 non-null   object  
 40  Asthma               10000 non-null   object  
 41  Services              10000 non-null   object  
 42  Initial_days          10000 non-null   float64 
 43  TotalCharge           10000 non-null   float64 
 44  Additional_charges    10000 non-null   float64 
 45  Item1                10000 non-null   int64  
 46  Item2                10000 non-null   int64  
 47  Item3                10000 non-null   int64  
 48  Item4                10000 non-null   int64  
 49  Item5                10000 non-null   int64  
 50  Item6                10000 non-null   int64  
 51  Item7                10000 non-null   int64  
 52  Item8                10000 non-null   int64  
 53  Soft_drink_numeric     10000 non-null   int64
```

```
dtypes: float64(11), int64(16), object(27)
memory usage: 4.1+ MB
```

There are no more NULL values in our dataset. We will now move on to outliers. We will obtain updated boxplots of the variables with outliers - Children and Income to see how they might have changed after imputation

```
In [35]: #Boxplot for Children
sns.boxplot(df, x='Children')
```

```
Out[35]: <Axes: xlabel='Children'>
```

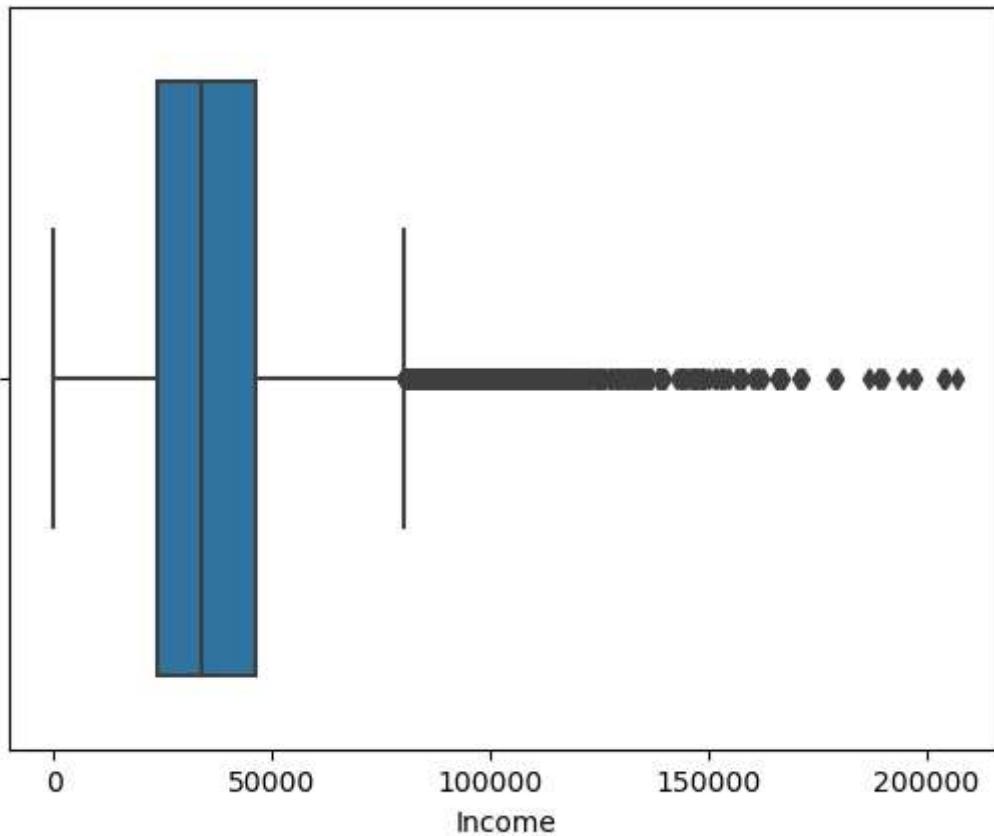


We can now see there appear to be 4 outliers in the number of children per patient - those patients who have 7-10 children are the outlier.

```
In [36]: #Obtain count of how many patients had >6 children
df['Children'][df['Children']>6].count()
```

```
Out[36]: 457
```

```
In [37]: #Box plot for Income
sns.boxplot(df, x='Income')
whisker_lines=plt.gca().lines
```



```
In [38]: #Obtain value of upper_whisker  
upper_whisker=whisker_lines[3].get_xdata()[0]  
print(upper_whisker)
```

80229.88

This shows that the outliers are patients who have an income greater than \$80,229

1. However, there are still too many outliers to be able to count with the boxplot, so we will need to use the count() function

```
In [39]: #Counting outliers in Income  
df['Income'][df['Income'] > 80229.88].count()
```

Out[39]: 705

The output shows us that there are 704 outliers, or patients with income greater than \$80229.88, in the Income column.

We will now review the statistics of Children and income

```
In [40]: df.Children.describe()
```

```
Out[40]: count    10000.000000
          mean     1.814000
          std      1.916969
          min      0.000000
          25%     1.000000
          50%     1.000000
          75%     3.000000
          max     10.000000
          Name: Children, dtype: float64
```

In [41]: `df.Income.describe()`

```
Out[41]: count    10000.000000
          mean    38872.450471
          std     25042.796229
          min     154.080000
          25%    23956.162500
          50%    33942.280000
          75%    46466.797500
          max    207249.130000
          Name: Income, dtype: float64
```

Determining outliers involves determining if the values are acceptable or not. While 10 children may seem like a lot, it is not unheard of. Additionally, \$207,249 income is a lot but also not unreasonable. At this time, without more information, I am choosing to retain all outliers in my dataset as I find the ranges to be reasonable.

We will now move on to apply Principal Component Analysis (PCA) to identify the significant features of the data set.

In [42]: `#import additional package for PCA
from sklearn.decomposition import PCA`

We will normalize all quantitative columns and store them in df_pca

In [43]: `#storing quantitative columns in df_pca using code from WGU course materials.
df_pca=df[['Population','Children','Age','Income','VitD_levels','Doc_visits','Full_meal']]`

In [44]: `# Normalize the columns by mean divided by standard deviation using code from WGU course materials.
df_pca_normalized = (df_pca-df_pca.mean())/df_pca.std()`

In [45]: `#PCA using code from WGU course materials.
pca=PCA(n_components=df_pca.shape[1])
print(pca)`

`PCA(n_components=11)`

We can see there are 11 principal components

In [46]: `#Normalizing PCA using code from WGU course materials.
pca.fit(df_pca_normalized)`

Out[46]: ▾ `PCA`

`PCA(n_components=11)`

```
In [47]: #PCA Normalization continued using code from WGU course materials.
df_pca=pd.DataFrame(pca.transform(df_pca_normalized), columns=['PC1','PC2','PC3','PC4'])

In [48]: #PCA Loadings using code from WGU course materials.
pca_loadings=pd.DataFrame(pca.components_.T, columns=['PC1','PC2','PC3','PC4','PC5','PC6','PC7','PC8'])

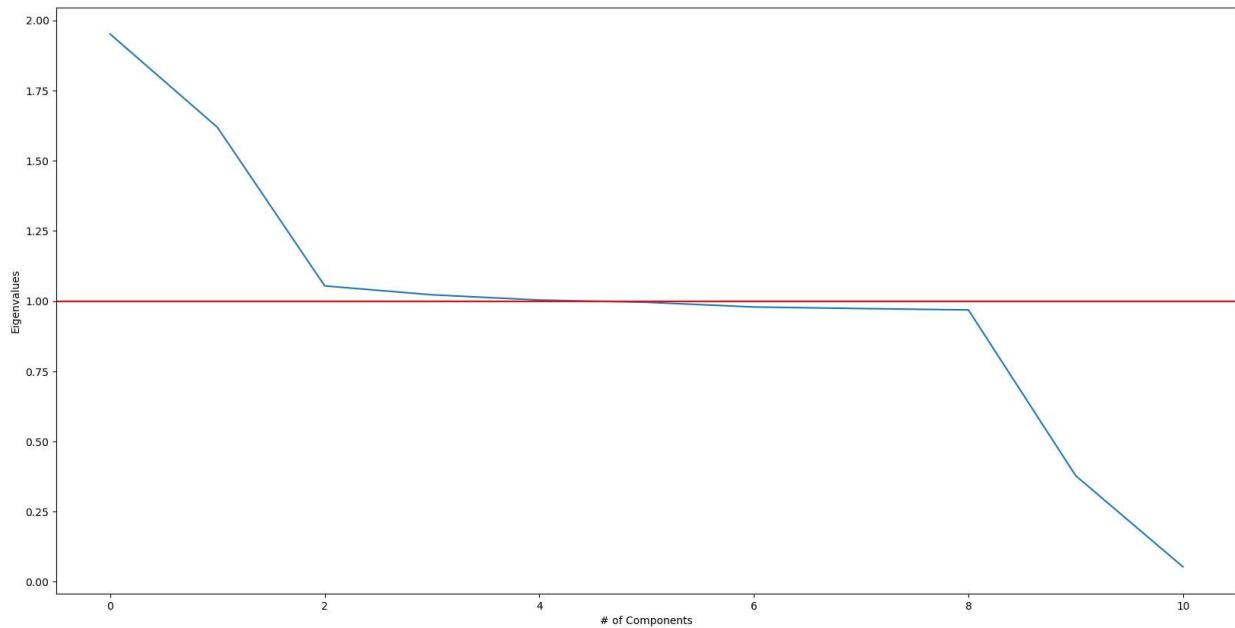
In [49]: #print Loadings using code from WGU course materials.
pca_loadings
```

Out[49]:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
Population	0.020664	-0.027012	0.503294	0.011609	-0.021267	0.545859	0.155000	0.2929
Children	0.004321	0.011337	0.143254	-0.055997	0.886951	0.143688	0.277812	-0.2468
Age	0.083046	0.700956	0.022109	-0.024584	-0.014235	0.004055	-0.015253	0.0219
Income	-0.006759	-0.005324	0.153030	0.616420	0.311495	-0.335102	-0.300044	0.5449
VitD_levels	0.540331	-0.052887	-0.291266	0.267675	-0.069787	0.088190	0.467171	0.1360
Doc_visits	-0.005253	0.012817	0.175189	0.626189	-0.179331	0.406596	-0.216702	-0.5299
Full_meals_eaten	-0.009220	0.036698	-0.552469	0.164700	0.231970	-0.008979	-0.227844	-0.2773
VitD_supp	0.033965	0.010677	0.424358	0.159839	-0.126966	-0.621372	0.382475	-0.3909
Initial_days	0.446473	-0.073504	0.316725	-0.315729	0.093103	-0.098161	-0.587243	-0.1618
TotalCharge	0.702174	-0.078309	-0.023420	0.003263	0.002287	0.005498	-0.016137	0.0025
Additional_charges	0.083681	0.701297	0.023997	-0.004434	-0.000742	0.011580	-0.003613	0.0212

In [50]: #obtaining eigenvalues using code from WGU course materials.
covariance_matrix = np.dot(df_pca_normalized.T, df_pca_normalized) / df_pca.shape[0]
eigenvalues = [np.dot(eigenvector.T, np.dot(covariance_matrix, eigenvector)) for eigenvector in eigenvectors]

In [51]: #plotting the eigenvalues using code from WGU course materials.
plt.figure(figsize = [20,10])
plt.plot(eigenvalues)
plt.xlabel('# of Components')
plt.ylabel('Eigenvalues')
plt.axhline(y=1, color='red')
plt.show()



```
In [52]: #list eigenvalue values using code from WGU course materials.  
eigenvalues
```

```
Out[52]: [1.9511903999693965,  
 1.6188378421815532,  
 1.0541224873896506,  
 1.0224182228569607,  
 1.0037755363342356,  
 0.9958246800154101,  
 0.9790089455874965,  
 0.9733461883648581,  
 0.9686663365082094,  
 0.3781461270225378,  
 0.05356323376974442]
```

This shows us that the first 5 Principal Components have a value greater than 1

```
In [53]: #confirm data set as cleaned  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 54 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        10000 non-null   int64  
 1   CaseOrder         10000 non-null   int64  
 2   Customer_id       10000 non-null   object  
 3   Interaction       10000 non-null   object  
 4   UID               10000 non-null   object  
 5   City              10000 non-null   object  
 6   State             10000 non-null   object  
 7   County            10000 non-null   object  
 8   Zip               10000 non-null   int64  
 9   Lat               10000 non-null   float64 
 10  Lng               10000 non-null   float64 
 11  Population        10000 non-null   int64  
 12  Area              10000 non-null   object  
 13  Timezone          10000 non-null   object  
 14  Job               10000 non-null   object  
 15  Children          10000 non-null   float64 
 16  Age               10000 non-null   float64 
 17  Education         10000 non-null   object  
 18  Employment        10000 non-null   object  
 19  Income             10000 non-null   float64 
 20  Marital            10000 non-null   object  
 21  Gender             10000 non-null   object  
 22  ReAdmis            10000 non-null   object  
 23  VitD_levels        10000 non-null   float64 
 24  Doc_visits         10000 non-null   int64  
 25  Full_meals_eaten   10000 non-null   int64  
 26  VitD_supp          10000 non-null   int64  
 27  Soft_drink          10000 non-null   object  
 28  Initial_admin       10000 non-null   object  
 29  HighBlood           10000 non-null   object  
 30  Stroke              10000 non-null   object  
 31  Complication_risk   10000 non-null   object  
 32  Overweight          10000 non-null   float64 
 33  Arthritis            10000 non-null   object  
 34  Diabetes             10000 non-null   object  
 35  Hyperlipidemia       10000 non-null   object  
 36  BackPain             10000 non-null   object  
 37  Anxiety              10000 non-null   float64 
 38  Allergic_rhinitis    10000 non-null   object  
 39  Reflux_esophagitis   10000 non-null   object  
 40  Asthma               10000 non-null   object  
 41  Services              10000 non-null   object  
 42  Initial_days          10000 non-null   float64 
 43  TotalCharge           10000 non-null   float64 
 44  Additional_charges    10000 non-null   float64 
 45  Item1                10000 non-null   int64  
 46  Item2                10000 non-null   int64  
 47  Item3                10000 non-null   int64  
 48  Item4                10000 non-null   int64  
 49  Item5                10000 non-null   int64  
 50  Item6                10000 non-null   int64  
 51  Item7                10000 non-null   int64  
 52  Item8                10000 non-null   int64  
 53  Soft_drink_numeric     10000 non-null   int64
```

```
dtypes: float64(11), int64(16), object(27)
memory usage: 4.1+ MB
```

```
In [54]: # Create copy of data set to perform remaining cleaning steps on
clean_df = df.copy(deep=True)
```

```
In [56]: clean_df.to_csv('C:/Users/Kmoik WGU/Desktop/KMoikclean_medical.csv') #This is showing
```

```
-----  
PermissionError Traceback (most recent call last)  
Cell In[56], line 1  
----> 1 clean_df.to_csv('C:/Users/Kmoik WGU/Desktop/KMoikclean_medical.csv')  
  
File ~\anaconda3\Lib\site-packages\pandas\core\generic.py:3772, in NDFrame.to_csv(self, path_or_buf, sep, na_rep, float_format, columns, header, index, index_label, mode, encoding, compression, quoting, quotechar, lineterminator, chunksize, date_format, doublequote, escapechar, decimal, errors, storage_options)  
    3761 df = self if isinstance(self, ABCDataFrame) else self.to_frame()  
    3763 formatter = DataFrameFormatter(  
    3764     frame=df,  
    3765     header=header,  
    (...)  
    3769     decimal=decimal,  
    3770 )  
-> 3772 return DataFrameRenderer(formatter).to_csv(  
    3773     path_or_buf,  
    3774     lineterminator=lineterminator,  
    3775     sep=sep,  
    3776     encoding=encoding,  
    3777     errors=errors,  
    3778     compression=compression,  
    3779     quoting=quoting,  
    3780     columns=columns,  
    3781     index_label=index_label,  
    3782     mode=mode,  
    3783     chunksize=chunksize,  
    3784     quotechar=quotechar,  
    3785     date_format=date_format,  
    3786     doublequote=doublequote,  
    3787     escapechar=escapechar,  
    3788     storage_options=storage_options,  
    3789 )  
  
File ~\anaconda3\Lib\site-packages\pandas\io\formats\format.py:1186, in DataFrameRenderer.to_csv(self, path_or_buf, encoding, sep, columns, index_label, mode, compression, quoting, quotechar, lineterminator, chunksize, date_format, doublequote, escapechar, errors, storage_options)  
    1165     created_buffer = False  
    1167 csv_formatter = CSVFormatter(  
    1168     path_or_buf=path_or_buf,  
    1169     lineterminator=lineterminator,  
    (...)  
    1184     formatter=self.fmt,  
    1185 )  
-> 1186 csv_formatter.save()  
1188 if created_buffer:  
1189     assert isinstance(path_or_buf, StringIO)  
  
File ~\anaconda3\Lib\site-packages\pandas\io\formats\csvs.py:240, in CSVFormatter.save(self)  
    236 """  
    237 Create the writer & save.  
    238 """  
    239 # apply compression and byte/text conversion  
--> 240 with get_handle(  
    241     self.filepath_or_buffer,  
    242     self.mode,  
    243     encoding=self.encoding,
```

```
244     errors=self.errors,
245     compression=self.compression,
246     storage_options=self.storage_options,
247 ) as handles:
248     # Note: self.encoding is irrelevant here
249     self.writer = csvlib.writer(
250         handles.handle,
251         lineterminator=self.lineterminator,
252         (...),
253         quotechar=self.quotechar,
254         )
255     self._save()

File ~\anaconda3\Lib\site-packages\pandas\io\common.py:859, in get_handle(path_or_bu
f, mode, encoding, compression, memory_map, is_text, errors, storage_options)
854 elif isinstance(handle, str):
855     # Check whether the filename is to be opened in binary mode.
856     # Binary mode does not support 'encoding' and 'newline'.
857     if ioargs.encoding and "b" not in ioargs.mode:
858         # Encoding
--> 859         handle = open(
860             handle,
861             ioargs.mode,
862             encoding=ioargs.encoding,
863             errors=errors,
864             newline="",
865             )
866     else:
867         # Binary mode
868         handle = open(handle, ioargs.mode)

PermissionError: [Errno 13] Permission denied: 'C:/Users/Kmoik WGU/Desktop/KMoikclean
_mcical.csv'
```

In []: